StreamReduce: A Data Stream Processor with MapReduce Parallelism

Ryan Lopopolo

Massachusetts Institute of Technology

`lopopolo@mit.edu`

May 15, 2012

1 Introduction

One of the fundamental problems MapReduce was created to solve is processing the growing amount of data in the world. MapReduce lends itself naturally to top-k-style computations, data sampling tasks, and other cases where perfect-fidelity results are not required. Now with the open source Apache Hadoop and Pig [11], it is easy to process and query the terabytes of data an organization may collect.

With terabytes of data, however, it is scarcely necessary to examine every entry to uncover trends, especially when one only cares about the most prominent trends or results. How does one decide which data to process then? Ultimately, this is problem dependent; however, selectively deciding which data to log is problematic because it can systemically mask some of the trends one would try to uncover through analysis. In this paper, I introduce StreamReduce, a distributed system for processing data streams. With StreamReduce I explore:

- Randomly discarding data while trying to maintain a target accuracy level.

- Processing data using a variety of com-

putation routines, some of which are less accurate than others.

Combining both of these techniques allows the cluster to save time by executing faster, save resources by avoiding the need to process every datum, and minimize the impact of faults in the execution of a job by treating failed tasks as discarded tasks.

Another limitation of MapReduce is batch processing. How does one incorporate new data into an already analyzed task? MapReduce simply says to rerun the job with the new data included. Others have solved this problem by batching data into chunks, continuously spawning MapReduce jobs to incorporate the new data, and incrementally storing updates in an external database instead of relying on the file system for enforcing the synchronization barrier.

A simpler way to incorporate new data is to treat data not as a constant, but as a stream. By processing data on demand, one avoids repeating work involved in reprocessing data, can save resources by spinning workers up and down as necessary, and ensures results are as fresh as possible by incorporating data into the result as soon as it arrives.

The goal of StreamReduce is to provide

a MapReduce-style framework for processing streams. Such a framework should include the ability to kill or discard tasks, be capable of providing workers with several different computation routines to choose among, and pass data between nodes using non-blocking routines.

2 Related Work

StreamReduce builds of previous work in the areas of MapReduce, Data Stream Management Systems, and approximation algorithms.

2.1 MapReduce

There is a rich research field around MapReduce. This section details work related to adapting MapReduce to a streaming environment.

The Hadoop Online Prototype [8] adds streaming to the MapReduce framework by removing the materialize phase between mappers and reducers and between phases of a MapReduce pipeline. Results are streamed from mappers to reducers as they become available. At any given point, each reducer maintains an in-progress snapshot of the MapReduce job's progress up to that point. This removes the synchronization barriers inherent in MapReduce and allows incremental progress to be observed.

This approach is in contrast to the continuous query approach taken by many other streaming systems. StreamReduce uses this snapshot approach.

2.2 Data Stream Management Systems

In their overview on data stream projects [6], a new class of data streaming applications called *Data Stream Management Systems* (*DSMS*es) is discussed. StreamReduce is a type of DSMS that includes some properties found in other systems.

The *Chronicle Data Model* [10] describes an append-only stream of tuples. Such streams lend themselves to incremental reads and incremental processing. In addition to incremental updates, the Chronicle Data Model also seeks to update results incrementally without storing the data itself.

The *Data Stream Model* [6] is a refinement of the Chronicle Data Model. It has the following properties: stream items may arrive out of order, streams are potentially unbounded in size, and elements in the stream are discarded after they are processed. This is the data model that StreamReduce uses.

The Aurora system [7] introduces the concept of *load shedding*, or compensating for resource overload by comparing job specific quality of service metrics with target levels and dropping items from the data stream. StreamReduce aims to incorporate load shedding to dynamically tune jobs in the face of impending time and resource constraints. On the other hand, it may be the case that parts of a processing pipeline struggle to provide enough data to keep the next stage busy. In this case, a streaming system would like to implement *early phase termination* [13], or preemptively kill a stage in the pipeline that is not capable of keeping all workers busy.

Aurora also introduces the non-blocking operators `filter`, `map`, and `union` for manipulating streams. StreamReduce implements these operators at the worker and fetcher nodes.

2.3 Approximation Algorithms

An alternative to load shedding to cope with high stream throughput is to perform approximate computations on stream data.

Dynamic Knobs [9] describes a mechanism to explore the accuracy–performance tradeoff space, dynamically adjusting program tuning parameters in response to changes in the execution environment. This allows programs to do things such as utilize more resources if they become available, use less power if power rates spike, or use less accurate computations in the face of an impending time constraint.

The most important piece Dynamic Knobs offers to the DSMS space is dynamic *peak load provisioning*. Their system can adjust downward the amount of resources, machines, etc. required for normal load and quickly adjust these levels back up to handle load spikes. Combined with Aurora's load shedding, peak load provisioning can be used to create a DSMS that uses less resources in both the normal and peak load states.

Accuracy-aware transformations [14] are a class of transformations that alter a computation according to a set of probabilistic weights and an accuracy target such that the computation performs more efficiently. StreamReduce uses the subclass of transformations called *substitution transformations* in worker nodes to adjust the runtime characteristics of per-datum computation.

The combination of load shedding, task killing and substitution transformations presents an opportunity for the results of a computation performed by StreamReduce to be distorted. Rinard [12] describes a method for bounding the distortion introduced when discarding tasks. By breaking a computation into tasks, which in StreamReduce consist of processing an individual datum at a worker, we can measure a *probabilistic distor-*

tion model for a job, which gives bounds on its accuracy at different task failure rates. I use this approach in the case study discussed in Section 5.

3 Design

StreamReduce is a distributed system for processing streams that can choose to kill tasks or limit the amount of work it performs per task. Its design is modular and there are five different types of nodes in a StreamReduce cluster:

Spout A node external to the cluster that is a source of streaming data.

Fetcher Communicates with spouts to provide data for workers.

Worker Processes data returned by fetchers. They are the source of parallel computation in the cluster.

Collector Aggregates results from all workers and makes the result accessible via an API.

Master Manages the cluster by bringing nodes online to complete jobs. It is responsible for all internode communication and storing to-be-processed data in several queues.

Figure 1 shows a system diagram of a StreamReduce cluster.

3.1 Spout

Spouts are network-connected machines that generate data for a StreamReduce cluster to process. Examples include a machine that tails a log file or an API like the Twitter tweet stream [**?**]. Fetchers depend on spouts as a source of streaming data.

Figure 1: Overview of data flow in the execution of a StreamReduce job. Load shedding occurs at point A. Task killing occurs at point B.

3.2 Fetcher

Fetchers are responsible for communicating with spouts and returning data for the workers to process. This level of indirection allows for a variety of different configurations: A spout can emit data at such a high rate that using multiple fetchers returns more data to the cluster; multiple fetchers can each talk to a different spout, for example tailing logs from multiple machines; or fetchers can be set up redundantly to ensure that every datum emitted by a spout is processed.

Fetchers may run indefinitely or return a status code indicating the fetching task is over.

3.3 Worker

Workers perform processing on data provided by fetchers. They are similar to mappers in the MapReduce framework. Each worker receives a fraction of the total input data and each worker's tasks are independent of other workers'.

Workers are responsible for fulfilling the main design goals of StreamReduce. Workers may choose to kill a task preemptively or, if it does not complete before a timeout, greed-

ily. Additionally, workers may choose among a variety of routines for performing computation on a datum. Routines are ranked on a $[0, 1]$ scale based on the estimated fraction of total work it requires. This rank is a proxy for the accuracy of a routine.

As it runs, a worker maintains an internal metric that estimates the fraction of work it has performed over all data it has seen. Workers use this metric when choosing which computation routine to run for a datum. A worker will choose the lowest ranked routine such that its estimated fraction of work is above a configured threshold. The total work estimator is used as a proxy for accuracy when a worker updates this metric. This metric also takes into account the number of tasks a worker has killed for any reason.

Because one of StreamReduce's goals is to process potentially infinite streams, workers do not save their progress to disk and are forgetful; once the master queries a worker for its progress and the worker successfully returns it, the worker clears all of its state and continues processing. As a consequence, workers are easy to add or remove from the worker pool because they don't need to be initialized with long-term job state. Addi-

tionally, as long as the master polls workers frequently enough, only a small fraction of the total work is lost upon worker failure. The master queries workers every 100ms by default.

3.4 Collector

Collector nodes aggregate the results of all workers. Each collector acts as a reducer would in the MapReduce framework if MapReduce used a single reducer for all of the mappers. Collectors differ from MapReduce reducers in that they receive their data piecewise, as it is produced from workers, as opposed to after all mappers have completed. This limits the types of stateless aggregation collectors can do to online operations. (Of course, a collector can choose to retain copies of individual worker results and re-reduce them each time more results are received, but this is not in the spirit of the framework.)

Collectors also expose an API endpoint for accessing the results of a job. This endpoint can be used by another spout to chain StreamReduce jobs together.

3.5 Master

There is one master node per StreamReduce cluster. It has a function similar to the jobtracker in MapReduce. New jobs are scheduled with the master and it initializes other nodes in the cluster with job state. For each job, the master schedules roles among nodes in a round robin manner to minimize the total number of outstanding roles on any one node.

The master is the only well-known node in the cluster. Other nodes in the network exclusively communicate with the master. The master and other nodes communicate with


Figure 2: The master communicates directly with all nodes to pass them control messages. Nodes respond directly to the master with their results.

each other by issuing HTTP requests to a simple server on each node. I decided to make use of a library [4] providing the server functionality because there was a low barrier to getting internode communication functioning. This design choice presents some problems and will be discussed further in Section 6.

Another consequence of the master being the only well-known node is that all data in the cluster must flow through it. The master maintains two queues per job for buffering job progress from fetchers, the data queue, and workers, the result queue, and periodically pushes this state out to workers and collectors, respectively.

Figure 2 shows how control messages are passed from the master to all other nodes in the cluster. The master communicates directly to each node and nodes reply directly to the master when they have completed their task. With these control messages, the master advances the state of the job. Such messages include telling a fetcher to fetch data from a spout, giving data to a worker to pro-

cess, and flushing results to collectors.

3.6 Configuring a Cluster for a Job

Apart from spouts, all nodes in a cluster modify their behavior based on the job they are running. The master must know the appropriate number of nodes to spin up; fetchers need to know how to talk to spouts; workers need to know how to compute results; and collectors need to know how to combine batched worker results.

StreamReduce achieves this configuration with a Jobfile. Jobfiles are instances of a special class provided by the StreamReduce framework. This class exposes several methods that can be invoked by a given node to provide configuration parameters or execute code on its behalf. Most of the routines that advance the progress of a StreamReduce job are wrappers around these Jobfile methods.

StreamReduce sends the source of the Jobfile to every node participating in the job and uses `eval` to obtain a local instance of the Jobfile class. An example Jobfile can be found at [2].

The use of `eval` to instantiate Jobfile objects and the implementation language's (ruby) ability to monkey-patch classes means that jobs can be modified on the fly just by sending an updated Jobfile to job participants. For example, a job can be made lower fidelity in the face of an impending time constraint, be made higher fidelity if more workers are added to the cluster, or gracefully kill all fetchers which allows the job to end after all of the various job queues are empty.

3.7 Fault Tolerance

StreamReduce can sustain the failure of any node other than the master. Because the master is responsible for shuffling data between fetchers, workers, and collectors, a job cannot make progress without the master.

Fetchers can fail. The only consequence of fetcher failure is that some of the input stream is lost forever. Because StreamReduce uses the Data Stream Model, it does not attempt to recover this data.

Worker failures may entail some lost work. Having the master poll each worker frequently for its progress mitigates the consequences of worker failure. For jobs that run for minutes or hours, 100ms of lost work is $\ll 0.1\%$ of the total. StreamReduce can simply treat these lost tasks as preemptively killed, which is already expected by the cluster.

Each collector is an exact replica of all others. Collector failure does not impact correctness of the results delivered by any other collector. If no collectors are online, the data meant for them sits in a queue on the master. Bringing another collector online will allow these and future results to be aggregated.

3.8 Killing Tasks and Working Less

A design goal of StreamReduce is to compute meaningful results from a stream and to do so while not processing every datum one-to-one. StreamReduce achieves this goal in two different ways. First, workers may kill tasks preemptively to avoid the time and resource costs associated with processing that task. Second, workers rely on the Jobfile to provide them with several different computation routines of differing fidelities that perform a fraction of the total work on a datum.

6

These behaviors appeal to the law of large numbers. It assumes that trends will still be visible in spite of looking at only a fraction of the data. This means that the computation routines and the weights assigned to them must be adapted to fit the distribution of the data before the job starts (or updated *in situ* via a new Jobfile).

4 Methodology

The cluster I used to test StreamReduce is a 5-node network consisting of 1 spout, 1 fetcher, 4 workers, 1 collector, and 1 master. Roles were partitioned among nodes as shown in Figure 3. Worker nodes were XVM [5] VMs provisioned with 128MB RAM each. The master, spout, fetcher, and collector ran on a mid-2010 Macbook Air laptop with 4GB of RAM. I configured the cluster to kill tasks on the master node, before they were shipped to workers.

The job I ran was a word counter. The collector emits a hash whose keys are words and whose values are the frequency the key occurred in the test corpus.

To run the job, I had to create a spout and find a source of data. For data, I acquired a Wikipedia dump of the Featured Articles category. I used Wikipedia's XML exporter [1] on January 9, 2012.

To serve the data, I transformed it into a YAML file containing a hash of

$$\texttt{article\_id} \rightarrow \texttt{fulltext}.$$

To store fulltext, I base64 encoded it and deflated it using zlib. This requires workers to inflate and decode the fulltext before they can process it further. The server reads from this YAML file and returns the resulting hash as JSON.


Figure 3: Topology of the test StreamReduce cluster used in the word count case study. Each square box represents a single machine. Each rounded box represents a StreamReduce role.

The fetcher communicates to the spout by requesting sequential article ids until it has fetched every article.

Workers are initialized with the three computation routines shown in Table 1.

The results of the word count job were then used in a script to compute the top 100 words. The cluster and word count job could be used in a two-stage top-k job. This script filters out SEO stop words [3] to achieve more interesting results.

I ran four trials of the word count job for every combination of work level and task kill rate I tested.

5 Evaluation

To evaluate StreamReduce, I explore the precision and recall of the word count test job

Weight	Description
1.0	Deflate and decode fulltext and count words in all of fulltext
0.25	Deflate and decode fulltext and count words in approximately a quarter of fulltext
0.1	Return the results of the last computation from a cache

Table 1: Computation routines used by StreamReduce workers.

as well as the speedup factor associated with varying the task kill rate and target work level.

5.1 Precision and Recall

For task killing to be an effective strategy in a top-k computation, the distribution of the top-k words must be stable. A stable distribution is one that is not altered by adding an additional article to or removing the last article from the corpus. Stable distributions are able to sustain the data lost in an approach like task killing.

Every data set has some base level of noise in it due to such things as how the data is sampled. StreamReduce introduces another source of noise by discarding data and substituting computation algorithms. If the noise introduced by StreamReduce is less than the data set's sensitivity to noise, the top-k words remain stable.

For example, if the distribution of words in the test corpus is close to uniform, the top-k words are bunched together and any perturbations in the tested data, for example, discarding an article that contains instances of the real top-k words, will affect the fidelity of the reported top-k.

The distribution of words in the Wikipeida data set is shown in Figure 4. The distri-


Figure 4: The frequencies of words in the Wikipedia Featured Articles corpus have a long-tailed distribution. This graph has had stop words [3] filtered out of it.

bution is long-tailed, which is a first step in verifying that this data set has a stable top k.

5.1.1 Precision

For the results given by StreamReduce's top-k computation to be meaningful, it is important for the distribution of the top-k words to be stable, which implies that task killing minimally impacts the head of the word distribution. Table 3 shows the top-10 words and their frequencies. Within the true top-10 words, the first word occurs twice as often as the tenth word, and the difference in frequencies of adjacently ranked words varies from 400–4000. Given that the data set has approximately 16 million words and 3500 articles, there are about 4500 words in any given article. Even the most frequently occuring words only appears on average 8 times per article. Gaps in the hundreds are stable in the addition or removal of tens of articles.

Figure 5 shows that even with a task kill rate of 0.8 and a work level of 0.4, meaning that the cluster only looks at 8% of the words in the corpus, the results are at least 80% ac-

Figure 5: Fraction of the top 100 results returned in the top 100 by the word count job.

Figure 6: Fraction of the bottom 10000 results returned in the result set by the word count job.

curate. Increasing the work level, which increases the fidelity of results emitted by workers, also increases the fidelity of the top 100 results.

5.1.2 Recall

Recall is sensitive to the stability of the distribution in the tail. Given that the Wikipedia data set is long-tailed, the tail of the distribution is very unstable. All of the words in the bottom-10000 have frequencies in the single digits. Because the distribution in the tail is so tightly-bunched, recall in the tail of the distribution is most dependent on the work level of the word count job. Less task killing and higher-fidelity worker computations benefit recall.

Anything less than perfect computation of a datum leaves out words with a frequency of one. Figure 6 shows how recall improves as the work level increases. Task killing has the same effect on recall as using low fidelity computation routines, so recall decreases as the task kill rate does. Task killing dominates the impact of work level on recall as demonstrated by the two lines converging in Figure 6.
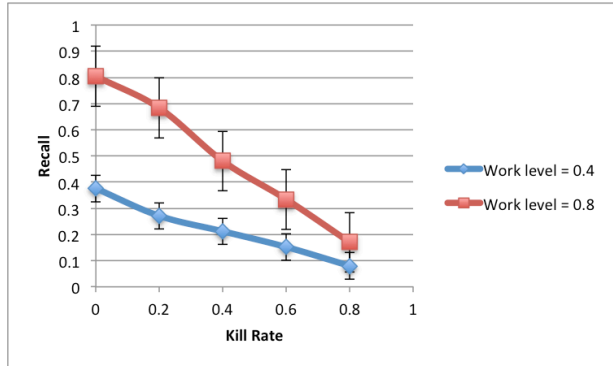
Poor recall statistics are acceptable be-

cause I am running a top-k job that focuses on good precision in the top 100. Because I use a weighted F-score, the poor recall's impact on measured accuracy is lessened.

5.1.3 Word Count Job Accuracy

Because I am running a top-k job, I compute the $F_{0.5}$-score of the StreamReduce results as opposed to the $F1$-score because I care more about precision of the top k than recall in the tail. Figure 7 shows that the $F_{0.5}$-score, or the weighted accuracy of the top-100 classifier, is always higher than the fraction of the input examined (task kill rate * work level). Setting the work level to 0.4 nets a greater increase in job accuracy versus fraction of words examined, with accuracy approximately 1.75 times higher than the fraction of words examined.

The average reduction in $F_{0.5}$-score caused by reducing the work level form 0.8 to 0.4 is 0.231 with a standard deviation of 0.025. I performed four additional experiments with the work level set to $\{0.8, 0.4\}$ and with the task kill rate set to $\{0.1, 0.7\}$. These experiments were not included in generating Figure 7. The average reduction in $F_{0.5}$-score between work level 0.8 and 0.4 in these exper-

9

Figure 7: $F_{0.5}$-score for the top-100 word count job. I measured the F-score with $\beta = 0.5$ because the goal of the job is to run a top-k filter on it, meaning I care more about precision than recall.

Task kill rate	Work level	Average execution time [s]	$F_{0.5}$-score
0.0	1.0	696	1.000
0.0	0.8	631	0.939
0.0	0.4	634	0.725
0.2	1.0	531	0.933
0.2	0.8	518	0.890
0.2	0.4	506	0.627
0.4	1.0	393	0.829
0.4	0.8	388	0.797
0.4	0.4	414	0.552
0.6	1.0	292	0.774
0.6	0.8	254	0.690
0.6	0.4	255	0.453
0.8	1.0	144	0.541
0.8	0.8	135	0.489
0.8	0.4	153	0.290

Table 2: Impact of task kill rate and work level on job execution time and $F_{0.5}$-score.

iments is 0.230 with a standard deviation of 0.056. These numbers show that a 0.4 reduction in the work level leads to approximately a 0.23 reduction in $F_{0.5}$-score.

5.2 Speedup

StreamReduce is able to achieve a speedup over a MapReduce-style computation through three mechanisms: providing multiple computation routines for the workers to choose from; varying the task kill rate; and varying the target work level, which influences the chosen frequency of computation routines per worker.

Running the perfect computation on one node took 1088 seconds. Compared to the fastest time in Table 2 using 4 workers, this is over a 700% speedup. There is much to be gained by processing data in parallel.

Table 2 shows the effect of varying the task kill rate and work level. The row with kill rate = 0.0 and work level = 1.0 represents the perfect, MapReduce-style, execution of the word count job. The only consistent reduction in execution time is a result of increasing the task kill rate. Increasing the task kill rate exhibits a linear reduction in execution time as the fraction of processed data linearly decreases.

Notably, simply by killing tasks, StreamReduce is able to achieve over a 5 times speedup compared to the perfect execution while still returning about 84 of the top 100 words.

5.3 Speed vs. Accuracy Trade-offs

Because the word distribution is so skewed toward the head, there is a great potential to speed up the calculation by killing tasks

10

| Work level = 1.0 | | Work level = 0.8 | | Work level = 0.4 | |
Task kill rate = 0.0		Task kill rate = 0.4		Task kill rate = 0.8	
33225	title	21184	title	5982	title
29356	web	18398	web	4689	web
23095	book	13591	journal	4225	journal
21821	journal	13429	book	3629	year
20124	year	12770	publisher	3617	publisher
19786	publisher	12550	url	3444	url
19478	url	12220	year	3410	book
17360	news	11511	news	3273	news
16685	accessdate	11082	united	2968	accessdate
16401	united	10691	accessdate	2582	world

Table 3: Example top 10 results and their frequencies from several different job configurations. The top 10 do not change much one changes the parameters to make the job less accurate.

without having to sacrifice too much accuracy. Table 3 demonstrates how little the top 10 changes in response to varying the job parameters. Additionally, Table 2 shows that the most significant time savings comes from increasing the kill rate and keeping the work level high yields about a 20 percentage point increase in $F_{0.5}$-score.

The best combination of parameters is a high work level and as high a task kill rate one can manage with respect to one's target distortion level.

6 Implementation Notes

This section describes deficiencies in the implementation of StreamReduce used to run the case study in this paper and goals for the next version of the system.

6.1 On-the-Fly Job Reconfiguration

The mechanism for altering execution of a job as it is running described in Section 3.6 is not yet implemented. Adding this feature to StreamReduce is a matter of adding an API endpoint to the servers run by each role and adding a bin-script to submit the new Jobfile.

6.2 Internode Communication

In Section 3.5 I mentioned design difficulties relating to the master and using HTTP requests for sending messages between nodes. I initially chose to use HTTP requests with POST payloads to send data and results between nodes because of its low barrier to entry; however, initiating many HTTP requests in succession is expensive. When killing jobs on worker nodes, the observed time savings was minimal because most of the time spent per datum involved shipping it from the master to the worker.

To mitigate the effects of this problem, I batched individual requests into groups of 5. This lowered the number of connections I was initiating but it also limited the parallelism of the cluster and removed workers from the pool for longer periods because they had more data to process at a time. A killed task on the

worker no longer freed it up to receive a new task immediately. Keeping workers removed from the pool caused the queues on the master to grow longer as the workers could no longer keep up with the rate the fetcher was producing data.

Eventually, I introduced a configuration option to kill tasks at the master. The master then decides to discard a task as it is popped from the queue, before it is ever sent to a worker. Doing so improved the execution time of the cluster but also meant that workers no longer had a full idea of how accurate they were being because they never killed any tasks themselves.

To address this problem, the master needs to keep a persistent channel open between it and all active nodes. Communication channels should be established when nodes come online. Possible solutions include writing to a socket or using a higher abstraction like AMPQ. AMPQ is appealing because it also relieves the master from handling queues.

6.3 Node Health

Currently, the master does not probe nodes to ensure that they are responsive. In the case that a node is serving a worker role for a job, this deficiency is masked. The node will simply not return the HTTP request sent to it to push it some data and the worker will remain removed from the pool. If the node is in the role of fetcher or collector, the master will stall.

All messages passed by the master to other nodes must be made non-blocking.

6.4 Job Management

There is no support for killing jobs or removing them from the master's data structures.

The only way to remove a job is to kill the master and all other nodes in the cluster.

7 Conclusion

I have extended the ideas of approximate computations to a MapReduce-style framework and shown that it is feasible to tune the input parameters of a job to achieve sufficient precision and recall for a particular application.

I have taken a different approach to streaming MapReduce by treating data not as a batch, but as a stream. At any given point in the execution of a job, the collectors will have accurate (so far as the job parameters allow) results for data seen until that point. This approach guarantees freshness and makes it easier to incorporate new data into an existing computation, yet it suffers from being limited to online aggregations.

The StreamReduce source code is open source and available at:

```
https://github.com/lopopolo/sr
```

Acknowledgments

References

[1] Export pages. `http://en.wikipedia.org/wiki/Special:Export`. Accessed January 12, 2012.

[2] lopopolo/sr. `https://github.com/lopopolo/sr`. Code repository for sr.

[3] Online seo guide: Google stop words - a comprehensive list of words google ignores. `http://www.link-assistant.com/seo-stop-words.html`. Accessed April 24, 2012.

[4] Sinatra. `http://sinatrarb.com`. Accessed May 4, 2012.

[5] XVM—virtual servers for MIT. `http://xvm.mit.edu/`. Accessed January 16, 2012.

[6] Brian Babcock, Shivnath Babu, Mayur Datar, Rajeev Motwani, and Jennifer Widom. Models and issues in data stream systems. In *Proceedings of the twenty-first ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*, PODS '02, pages 1–16, New York, NY, USA, 2002. ACM.

[7] Don Carney, Uğur Çetintemel, Mitch Cherniack, Christian Convey, Sangdon Lee, Greg Seidman, Michael Stonebraker, Nesime Tatbul, and Stan Zdonik. Monitoring streams: a new class of data management applications. In *Proceedings of the 28th international conference on Very Large Data Bases*, VLDB '02, pages 215–226. VLDB Endowment, 2002.

[8] Tyson Condie, Neil Conway, Peter Alvaro, Joseph M. Hellerstein, Khaled Elmeleegy, and Russell Sears. Mapreduce online. Technical Report UCB/EECS-2009-136, EECS Department, University of California, Berkeley, Oct 2009.

[9] Henry Hoffmann, Stelios Sidiroglou, Michael Carbin, Sasa Misailovic, Anant Agarwal, and Martin Rinard. Dynamic knobs for responsive power-aware computing. In *Proceedings of the sixteenth international conference on Architectural support for programming languages and operating systems*, ASPLOS '11, pages 199–212, New York, NY, USA, 2011. ACM.

[10] H. V. Jagadish, Inderpal Singh Mumick, and Abraham Silberschatz. View maintenance issues for the chronicle data model (extended abstract). In *Proceedings of the fourteenth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems*, PODS '95, pages 113–124, New York, NY, USA, 1995. ACM.

[11] Christopher Olston, Benjamin Reed, Utkarsh Srivastava, Ravi Kumar, and Andrew Tomkins. Pig latin: a not-so-foreign language for data processing. In *Proceedings of the 2008 ACM SIGMOD international conference on Management of data*, SIGMOD '08, pages 1099–1110, New York, NY, USA, 2008. ACM.

[12] Martin Rinard. Probabilistic accuracy bounds for fault-tolerant computations that discard tasks. In *Proceedings of the 20th annual international conference on Supercomputing*, ICS '06, pages 324–334, New York, NY, USA, 2006. ACM.

[13] Martin C. Rinard. Using early phase termination to eliminate load imbalances

at barrier synchronization points. In *Proceedings of the 22nd annual ACM SIGPLAN conference on Object-oriented programming systems and applications*, OOPSLA '07, pages 369–386, New York, NY, USA, 2007. ACM.

[14] Zeyuan Allen Zhu, Sasa Misailovic, Jonathan A. Kelner, and Martin Rinard. Randomized accuracy-aware program transformations for efficient approximate computations. In *Proceedings of the 39th annual ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '12, pages 441–454, New York, NY, USA, 2012. ACM.